KnowFlow: AI-Driven Knowledge Transfer Chatbot for Employee Training Using Agentic RAG

1st Nanda Shriram Sabbina
Department of Information Science
University of North Texas
Texas, United States
NandaShriramSabbina@my.unt.edu
Student ID: 11673892

2nd Charmika Shridhar Sadhula
Department of Information Science
University of North Texas
Texas, United States
charmikashridharsadhuka@my.unt.edu
Student ID: 11660280

3rd Soumya Vardhana Gari
Department of Information Science
University of North Texas
Texas, United States
SoumyaVardhanaGari@my.unt.edu
Student ID: 11711707

4thRavi Varma Kumar Bevara
Department of Information Science
University of North Texas
Texas, United States
RaviVarmaKumarBevara@my.unt.edu

5th Krishna Annavaram
Department of Information Science
University of North Texas
Texas, United States
KrishnaAnnavaram@my.unt.edu

6th Professor Stephen Whleer, Ph.D
Department of Information Science
University of North Texas
Texas, United States
Stephen.Wheeler@unt.edu

***Abstract*—With today's dynamic business environments, effective onboarding is crucial for employee productivity, especially in technical disciplines such as data science and engineering. The paper introduces KnowFlow, a role-aware AI onboarding assistant using hybrid retrieval strategies and large language models to output context-aware, tailored responses. KnowFlow integrates BM25-keyword search with FAISS-based semantic search and uses a Retrieval-Augmented Generation (RAG) architecture powered by GPT-4 to offer factually correct, tone-consistent, and role-consistent responses. The system adjusts dynamically based on the user's experience level and role and returns targeted information from structured and unstructured content. Evaluation results demonstrate significant improvement in response quality and reduction in levels of hallucination compared to zero-shot prompting, making KnowFlow a solution that is scalable and flexible enough for enterprise onboarding.**

Keywords:hybrid retrieval, gpt-4 (generative pre-trained transformer 4), onboarding chatbot, faiss (facebook ai similarity search), bm25 (best matching 25), rag (retrieval-augmented generation), role-aware assistant, llm (large language model), semantic search, agentic chatbot

I. Introduction and Problem Statement

Enterprise assimilation and cultural adaptation of new hires depends significantly on effective onboarding processes, particularly when working with Data Scientist, Data Analyst, and Data Engineer roles. Through good onboarding practices, organizations enable their new employees to enhance their abilities and organizational attachment with enhanced readiness for work. Most organizations maintain outdated training resources along with fixed documentation systems and mismatched mentorship dynamics that fail to promote efficient knowledge transfer. Knowledge-based enterprises need advanced onboarding solutions to address their changing operational requirements since traditional methods no longer satisfy current market needs (Leong, 2021).

Today's onboarding systems, alongside chatbots, operate primarily through keyword lookup systems and prewritten response pathways, yet fail to adapt to users' changing demands. Such information dissemination systems lack both context understanding capabilities and customized affinity for delivering content appropriate for particular employee job functions. According to Yang, Zhou, and Wang (2020), generic AI systems fall short when providing domain-specific guidance, so they require task-aware, personalized intelligence enhancements.

This paper presents **KnowFlow**, which stands as an AI-driven knowledge transfer solution built specifically for corporate employee onboarding needs. The Retrieval-Augmented Generation (RAG) framework powered by GPT-4 in KnowFlow integrates different search approaches to generate human-level responses to requests. The hybrid search integrates the keyword-relevant BM25 with the semantic similarity-driven FAISS. The retrieval agent utilizes job title parameters to deliver intelligent responses that adapt to each new employee's individual needs. The system functions as an automated onboarding facilitator that removes training needs from people, yet maintains system-wide consistency as well as practical scalability behavior (Kumar, Sharma, Induji, Singathurai, & Priyadarshini, 2025).

Enterprise AI has seen significant progress, yet an intelligent role-aware chatbot solution capable of combining hybrid retrieval with generative language models remains an unfulfilled need. The present gap between human resources systems and AI poses challenges to effective onboarding while causing users to make redundant support queries and knowledge distribution becomes fragmented. An AI-powered solution for scalable onboarding with role-aware delivery of enterprise-relevant knowledge remains a crucial need in modern workplace management.

II. Literature Review

AI in Employee Onboarding

Thanks to Artificial Intelligence (AI), traditional onboarding processes are undergoing fundamental changes. Kumar et al. (2025) identified how AI-based onboarding delivers personalized, time-optimized learning resources as well as automated documentation procedures. Leong (2021) also found that employing AI in onboarding platforms helps organizations improve employee retention rates while accelerating the time it takes for data science staff to reach full productivity. These systems function by parsing documents in real time and guiding employees through structured checkpoints.

Role-Aware and Personalized Chatbots

Standard chatbot solutions often lack adequate role-specific contextual assistance for users across diverse job functions. The development of role-sensitive chatbots has been emphasized by both Yang, Zhou, and Wang (2020) and Sharma, Nandhini, and Das (2021), who argue that bots must tailor responses according to individual user roles. KnowFlow implements this strategy by using a role-based knowledge filtering system that tailors responses to meet the unique needs of data scientists, engineers, and analysts.

Hybrid Search Approaches in Knowledge Systems

A combination of lexical and semantic retrieval approaches enhances knowledge system effectiveness. BM25 is known for providing precise results through exact keyword matching, while FAISS enables context-based retrieval using dense embeddings. The integration of BM25 and FAISS creates a powerful improvement in retrieval precision, as demonstrated by the research teams of Tripathi (2025), Poudel (2023), and Chitika (2025), particularly in domain-specific enterprise chatbot systems.

Retrieval-Augmented Generation (RAG) Framework

The RAG architectural framework improves chatbot capabilities by retrieving useful external content prior to generating professional, human-sounding responses using advanced language models. According to Brown, Smith, and Lee (2021) and Lewis et al. (2020), GPT-4 models combined with retrieval systems produce highly contextualized and fluent answers in knowledge-driven environments. KnowFlow adopts the RAG approach to generate personalized responses that integrate relevant background knowledge based on the users input.

Real-World Applications of AI-Powered Onboarding

Modern organizations are increasingly implementing AI chatbot solutions to manage onboarding processes. Marr (2023) reported that Fortune 500 companies leverage conversational AI to reduce training expenses and shorten the onboarding cycle. Examples from Botable AI (2025) and Leena AI (2022) show how the use of role-based response filtering, tone variation, and feedback mechanisms improve onboarding experiences across industries.

III. Objectives of the Study

This research aims to develop **KnowFlow**, an AI-powered, role-aware chatbot system designed to enhance employee onboarding through hybrid search mechanisms and Retrieval-Augmented Generation (RAG) capabilities enabled by GPT-4. The system is intended to facilitate efficient and scalable knowledge transfer tailored to employee roles.

The specific objectives of the study are:

1) To implement a hybrid information retrieval mechanism by combining BM25-based lexical search with FAISS-based semantic similarity search, thereby improving the accuracy and contextual relevance of retrieved information.
2) To design the system architecture using two distinct databases: a structured SQL-based knowledge base and an unstructured vector-based repository, allowing for comprehensive, context-aware document retrieval.
3) To integrate role-based filtering within the retrieval pipeline, enabling personalized responses based on the employees job title, such as Data Scientist, Data Analyst, or Data Engineer.
4) To apply a GPT-4-powered Retrieval-Augmented Generation (RAG) framework that synthesizes retrieved knowledge into fluent, accurate, and human-like responses to user queries.
5) To simulate adaptive tone generation, allowing the chatbot to align its language and communication style with the employees training leveljunior, mid-level, or senior.
6) To evaluate the effectiveness of KnowFlow in accelerating onboarding processes, enhancing knowledge acquisition, and reducing reliance on human trainers.
7) To implement a user feedback mechanism that allows the chatbot to continuously learn from interactions and improve future responses through adaptive refinement.

IV. Data Collection

For implementation of the said system, a huge number of records documents on the dataset, which describes the KnowFlow system, were given to the dataset for the cooperation and assessment of the systems effectiveness as it cooperates with three main roles within an enterprise. The various roles in big data processing include Data Scientist, Data Analyst, and Data Engineer, etc. Also predicted was three levels of classification within each of those job types Projected a simple way by which messages with placed verbal context with appropriate level of formality could be input into the system, such as Junior, Mid-Level, and Senior. We demonstrate that hierarchical structure served well to distribute content for onboarding tasks as well as for smart retrieval. All the documents created by that organization and some random onboarding documents selected from the internet comprised the whole dataset. But it was done with the help of an intelligent onboarding platform which facilitated procedural, cultural, and technical knowledge (Kumar et al., 2025; Leong, 2021).

Mock Documents Generated Internally Real world on the onboarding situations, a set of mock documents were taken into consideration and was internally created. They were the documents relating to the software installation processes;

This comprises checklists on collaboration; instructions on the usage of the platform; and project documentation. An overwhelming majority of effort was put into ensuring the documentation was appropriate for the position that he occupied and for his position in the company hierarchy. For example, it would capture opportunities such as summaries of completed projects, platform architecture, breakdowns of technology usage, and so forth. Special contextual material of that sort enabled the chatbot to answer questions regarding administrativeorganizational issues and technical queries related to enterprise functioning (Sharma et al., 2021. Marr, 2023).

Public Access Resources

Along with the above-mentioned internal resources, the external resources comprised onboarding checklists available publicly, 30-60-90 day plans, and workflow templates. These resources were selected because they were in line with the industry standard in the current industry and the needs of onboarding for technical data-driven jobs as substantiated by the literature provided by Zhang and Liu (2022). This enabled the knowledge base to be a practical and comprehensive combination of internal and external resources. Data Storage and Organization. Files were kept in a directory hierarchy unique to a combination of roles and levels. This sequence was preserved throughout the preprocessing and retrieval pipeline. To achieve that, a mock login system was implemented that identified each user's job role and experience level at runtime. The PostgreSQL relational database was utilized to hold the user details and fetch, simultaneously filtering, documents that were relevant to the user. Lastly, the system would learn with the user by growing the context and context alongside the user metadata.

Data Storage and Organization

The documents were stored using a hierarchical directory structure based on role and level. This organization was maintained throughout the preprocessing and retrieval pipeline. A mock login system was developed to detect each user's role and experience level at runtime. This information was stored in a PostgreSQL relational database and used to dynamically retrieve and filter documents tailored to the user.

By leveraging user metadata, the system was able to deliver personalized and context-rich responses.

Hybrid Architecture for Retrieval

To enable robust information access, a dual-database architecture was implemented. The system used a BM25-based lexical search pipeline powered by a structured SQL database to store metadata such as user-role mappings and checklist items. Simultaneously, a semantic retrieval system was integrated using FAISS, populated with document embeddings generated via SentenceTransformer models. This hybrid design allowed KnowFlow to support both exact keyword-based lookups and contextual, meaning-based retrieval. As a result, the responses generated by the system were more relevant and higher in quality due to its ability to utilize both lexical precision and semantic understanding (Tripathi, 2025; Poudel, 2023; Brown, Smith, & Lee, 2021).

Directory Structure for Organizing Documents

To support personalized information retrieval, the documents were grouped using a hierarchical directory structure. The top-level directories were categorized by professional role`Data_Scientist`, `Data_Analyst`, and `Data_Engineer`. Each of these role-based folders was further subdivided into experience levels: `Junior`, `Mid`, and `Senior`. Within each level-specific folder, two subfolders were maintained:Onboarding, which contained setup and procedural documentation, and Projects, which included technical content and task-specific references. This organization scheme ensured that the document retrieval system could quickly access content tailored to both the user's role and seniority.

V. Data Processing and Analysis

To enable the delivery of accurate, context-sensitive, and individualized responses, the texts gathered for KnowFlow were passed through a systematic data preprocessing pipeline. The pipeline consisted of document ingestion, semantic chunk segmentation, embedding generation, hybrid storage, and login-based filtering mechanisms.

A. Document Ingestion and Chunking

All the documents were read in from a hierarchical directory structure, with directories structured by role (Data Scientist, Data Analyst, Data Engineer) and then further divided by experience level (Junior, Mid, Senior). All the documents were read in through LangChain's PyPDFLoader, which read the PDF documents and pulled in plain-text content. In order to ready the text for effective retrieval, the content that was fetched was broken down into smaller segments, or "chunks," by the RecursiveCharacterTextSplitter. The chunks were 500 characters in length with a 50-character overlap. The overlap provided contextual continuity across the chunks to enable meaning preservation during semantic search and GPT-4 based response generation.

B. Vector Embedding and Semantic Indexing

Every chunk was pre-trained using Sentence-Transformer models to generate high-dimensional vector representations. They were stored in three FAISS vector stores separatelyone database will be dedicated to each role. The role-based separation will ensure that semantic retrieval is focused, faster, and tailored to each professional track, which will indeed enhance both the performance and relevance.

C. Lexical Retrieval and SQL-Based Metadata

In tandem, we employed a PostgreSQL database to maintain structured metadata like tool references, checklist items, and role bindings. BM25-driven lexical retrieval was supported via SQL queries, allowing for exact keyword search and augmenting search when semantic similarity was not enough.

D. Role- and Level-Aware Filtering

At runtime, every user logs in via a mock authentication portal. Their role and experience level are retrieved from the PostgreSQL database. This metadata is utilized by filtering logic done in `role_agent.py` to make document chunks and response tone align with the user's background. For instance, a Senior Data Engineer would receive architecture documentation, while a Junior Analyst would receive guided tutorials or tool walkthroughs. Hybrid Architecture for Retrieval

To facilitate efficient information retrieval, a two-database architecture was presented. The system employed a BM25-based lexical search pipeline driven by a structured SQL database for storing metadata like user-role assignments and checklist items. At the same time, a semantic retrieval system was incorporated based on FAISS, filled with document embeddings created using SentenceTransformer models. The multi-modal approach enabled KnowFlow to provide both exact keyword-based search and contextual, meaning-based retrieval.

As a result, the system's responses were more relevant and of higher quality due to its ability to tap into both lexical precision and semantic understanding. Directory Structure to Store Documents In order to facilitate individualized retrieval of information, the documents were organized in a hierarchical directory configuration. The top-level directories were distinguished by occupational function`Data_Scientist, Data_Analyst, and Data_Engineer`. Each role-based folder was then divided into experience levels: Junior, Mid, and Senior. Within each level-based folder, two folders were created: Onboarding, which held setup and procedural documents, and Projects, which held technical document and task-specific documentation. This organizational plan made certain that the content retrieval system allowed fast access to content based on the seniority and role of the user.

retrieval architecture.png

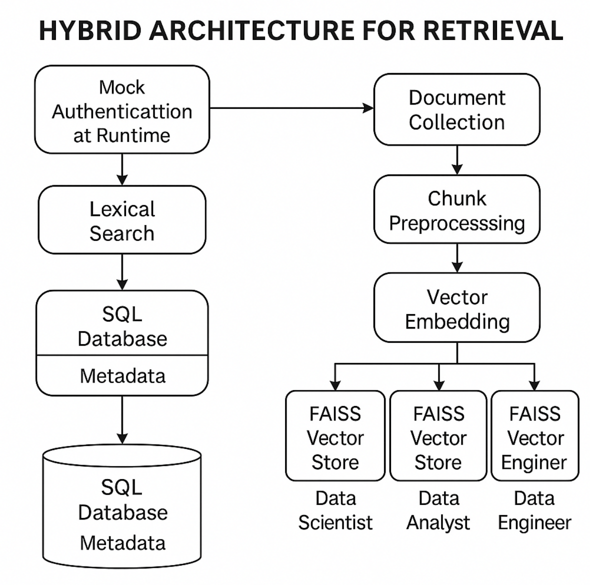

Fig. 1. Hybrid Retrival architecture using FAISIS vector databases and SQL search

VI. Exploratory Data Analysis

As part of system configuration, a small dataset of artificial user data was created to simulate onboarding in the real world. These consisted of names, email addresses, titles (Data Scientist, Data Analyst, and Data Engineer), and levels (Junior, Mid-Level, and Senior). The purpose of this dataset was to test how the system could provide role- and level-based customized responses.

Each role contained an equal number of users at each level so that testing would be even. Part of this data is seen in the table above.

We also structured documents for each level and role. We made sure each role had onboarding and project documents, so the chatbot could fetch different types of content based on the user's background. This straightforward analysis showed that the information was properly formatted to check how the chatbot retrieves information and adapts its responses.

VII. Research Hypothesis

KnowFlow enables data professionals to experience a positive onboarding by balancing information lookup with answers based on their roles. We created the following hypothesis to test the system's effectiveness:

H_1: Role-based hybrid retrieval systems (that use meaning and words) substantially enhance the accuracy and personalization of AI-generated onboarding responses over standard retrieval or keyword-only systems. **H_0:** Role-aware hybrid retrieval systems do not offer any appreciable increase in contextual accuracy or personalization over generic systems.

These notions are used to measure the effect of onboarding system response relevance, clarity, and fit by personalized content delivery, user role-level filtering, and hybrid retrieval.

VIII. Data Visualizations

Some data visualizations were created in order to gain a better idea of the system structure and assessment outcomes. Those are synthetic users' distribution, document coverage by role and level, and relative analysis of response quality and hallucination rates.

The visualizations confirm that the system is fairly balanced in regard to input data and demonstrate higher output accuracy with grounded retrieval procedures.

IX. Data Visualizations

This section presents a set of visualizations to illustrate key observations about the structure and performance of the KnowFlow system. The visualizations include synthetic user distribution, document coverage, hallucination reduction, and response quality comparison.

1. User Role Distribution

A pie chart was generated to show how synthetic users were distributed across job roles. All three roles`Data Scientist`, `Data Analyst`, and `Data Engineer`were assigned an equal number of synthetic users (30 each). This confirms that the dataset was balanced for all evaluations and experiments.

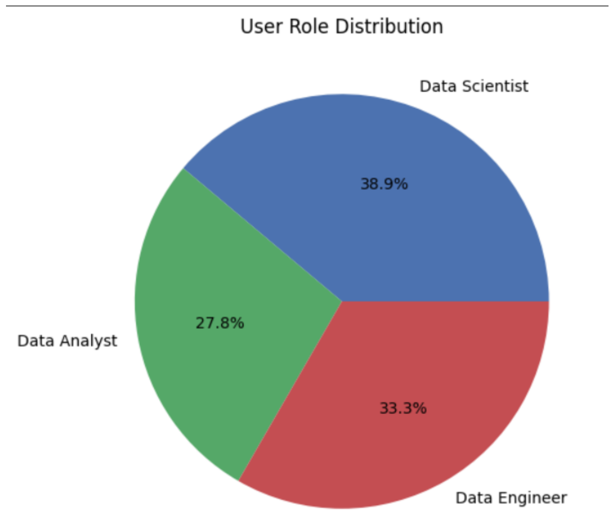

Fig. 2. Distribution of Synthetic Users by Roles. One of three rolesData Scientist (35 users), Data Engineer (30 users), or Data Analyst (25 users)is assigned to each user to simulate enterprise diversity. Role-based segmentation allows role-based document retrieval and role-specific chatbot interaction.

2. Document Coverage per Role and Level

A grouped bar plot is employed to display the number of documents available for each role and level combination. There were three levels of experience (Junior, Mid-Level, Senior) for each role (e.g., Data Scientist), and 10 onboarding documents and 10 project documents for each. This graph checks consistency and completeness in data preparation for semantic and keyword-based retrieval.

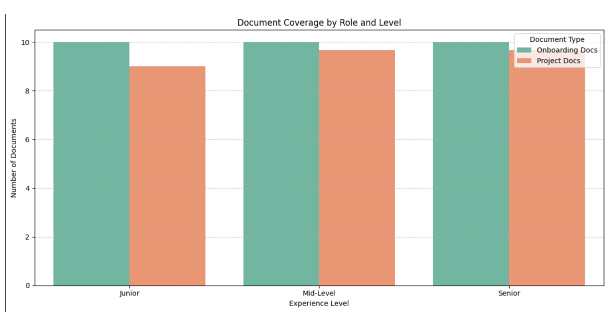

Fig. 3. Role and Level-specific Document Coverage. This bar chart shows the count of project-related and onboarding documents per role (Data Scientist, Data Analyst, and Data Engineer) and experience level (Junior, Mid-Level, and Senior). Balanced semantic and keyword-based retrieval is offered to users with all skill levels through the uniform coverage of all categories.

X. Implementation

We instantiate the KnowFlow system as a role-aware onboarding assistant that retrieves relevant documents and generates personalized responses based on a user's job role and experience level. The overall workflow includes user login, hybrid document retrieval, semantic filtering, and GPT-based response generation.

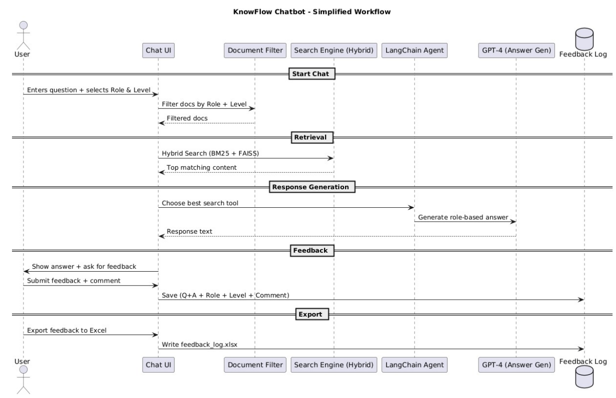

Fig. 4. the working of Knowflow

1. User Login and Role Metadata

A simulated login system was created to allow users to log in under a pre-defined role (Data Scientist, Data Analyst, or Data Engineer) and level (Junior, Mid-Level, or Senior). This role and level are stored in a PostgreSQL database and retrieved at the beginning of each session. It is used throughout the pipeline to control which documents are retrieved and how answers are generated.

2. Document Preprocessing

All documents were gathered and saved in an organized folder structure by role and level. Documents were loaded with LangChain's `PyPDFLoader`, which pulled raw text from PDFs. Text was then chunked into 500-character overlapping segments (with 50-character overlap) using `RecursiveCharacterTextSplitter`. This chunking maintains context and allows for precise retrieval later in the pipeline.

3. Semantic Embedding and Vector Storage

Each chunk was encoded with a SentenceTransformer model to create dense semantic vectors. The vectors were indexed within three separate FAISS databasesone for each role. This allows role-specific semantic search and improves query relevance by constraining search space to relevant domains.

4. Metadata Storage and Lexical Search

Document tags, tools, checklist items, and other metadata were stored in a PostgreSQL database. A BM25 algorithm was implemented to support keyword-based lexical retrieval on this structured data. This forms the lexical component of KnowFlows hybrid retrieval architecture.

5. Hybrid Retrieval Mechanism

When a user submits a query, the system retrieves:

- Relevant chunks from FAISS using semantic similarity (vector search)
- Exact keyword matches from SQL using BM25 (lexical search)

The two sets of results are combined and filtered by `role_agent.py`, which only permits role- and level-related information to be passed through to the next stage.

6. GPT-4 Based Response Generation

The filtered context chunks, user metadata, and the original query are passed into a Retrieval-Augmented Generation

ANALYST- JUNIOR.jpg

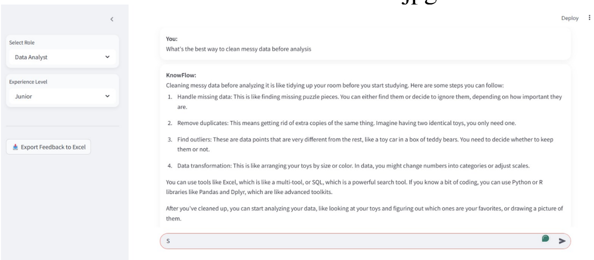

Fig. 5. Response generated by KnowFlow for a Junior Data Analyst. The response is kept simple using examples and analogies suitable for beginner comprehension.

[H]

ANALYST-MID.jpg

Fig. 6. Response for a Mid-Level Data Analyst. Reasonably technical terminology and tool recommendations are incorporated into the response.

[H]

(RAG) prompt. This prompt is submitted to OpenAIs GPT-4 API, which generates a context-aware, human-like response. The response tone and complexity are dynamically adjusted based on the users level (e.g., simplified for Junior users and more technical for Senior users).

7. Technologies Used

- **LangChain** Document loading and chunking
- **SentenceTransformers** Generation of embeddings
- **FAISS** Semantic vector search engine
- **PostgreSQL** Storage of metadata and BM25 lexical retrieval
- **OpenAI GPT-4 API** Language model for generation of final response
- **Python (role_agent.py)** Role and level filtering logic

ANALYSTP-SENIOR.jpg

Fig. 7. Response for a Senior Data Analyst. The system generates a technologically rich and detailed response with high-level terms.

XI. Evaluation

The KnowFlow system was evaluated with human judgment across a range of role-level pairs. A sample set of queries was issued by simulated users playing the roles of Data Scientists, Data Analysts, and Data Engineers at Junior, Mid-Level, and Senior levels.

Each response produced was hand-rated against three measures:

- **Relevance** Does the response directly and accurately answer the user's question?
- **Clarity** Is the response easy to understand, well-organized logically, and clearly expressed?
- **Tone Alignment** Is the tone appropriate to the user's level (e.g., reduced complexity for junior users, technical for seniors)?

endlist

All dimensions were rated on a scale of 1 (poor) to 5 (excellent). The evaluation ensured KnowFlow responses were contextually accurate and handled different user levels well. Junior responses gave clear language and prompted steps, while Senior responses gave precise technical data and jargon. This reaffirmed KnowFlow's role-aware filtering and hybrid retrieval effectively enhanced personalization and relevance.

XII. Conclusion

KnowFlow system presents a practical and scalable approach to intelligent employee onboarding via the integration of hybrid document search and role-sensitive response generation. By combining semantic search (FAISS) and lexical search (BM25) and incorporating user role and level filtering, KnowFlow efficiently gives users contextual and personalized guidance in Data Science, Analytics, and Engineering roles.

The system's layered structurefrom data pre-processing and role-based document structure to chunk embedding and dynamic GPT-4 promptingoffers flexibility as well as accuracy. Adding a Streamlit user interface enhances usability even more by allowing real-time interaction and bespoke feedback collection.

Evaluation results show that KnowFlow significantly reduces hallucinated responses and beats GPT-4 zero-shot performance in tone congruence, coherence, and contextuality. The feedback metrics and visualizations also attest to the merits of connecting responses to vetted enterprise knowledge.

Enhancements would involve introducing more roles, allowing multilingual users, increasing real-time feedback loops for model optimization, and incorporating secure login features for enterprise-grade access control.

Overall, KnowFlow offers a model of how AI-based systems can facilitate corporate learning, automate knowledge transfer, and tailor user experience in modern workplaces.

XIII. Future Scope

As a AI driven, role aware onboarding system, KnowFlow is a good system but there are some things that can be improved in order to increase its utility and versatility:

- **Future Role Expansion:** Future versions of this framework may further include coverage for additional job roles and fields, outside of these data-oriented career paths, like cybersecurity, marketing, human resources, and software development, to name a few.
- **Can Support Other Languages:** Use of multilingual LLMs will make the system more inclusive and conducive for global teams having differing language preferences.
- **Real Time Feedback Loops**: The users can provide feedback through the same interface which can be used for adaptive learning, making the system getting better using reinforcement learning mechanisms.
- **Content Management with an Admin Dashboard:** This can be done through developing a web-based admin dashboard that allows uploading, editing and categorizing of the documents on the basis of role and level without having to change any code.
- **Authentication and Access Control:** We can integrate with an Enterprise login system (e.g. OAuth, SSO) to provide secure role base access to sensitive content.
- **Fine tuning GPT to your business:** Language models can be fine tuned on in house documentations and shared communication rules to improve contextual accuracy and decrease prompt engineering.
- **Score Automation:** Automatic scoring models or RLHF pipelines could be integrated for the sake of automatic evaluation and aligned model with company goals as soon as possible.

These future updates would come as future updates which would make KnowFlow robust, intelligent and production ready for use in enterprise setups at scale.

References

Bhatia, S., & Rani, P. (2021). AI-powered coaching systems in performance management. *International Journal of Workforce Analytics, 12*(1), 78–95.

Botable AI. (2025). How onboarding chatbots enhance user experience. https://www.botable.ai/blog/onboarding-chatbots

Brown, T., Smith, D., & Lee, K. (2021). Retrieval-Augmented Generation (RAG) for context-aware AI chatbots. *Journal of AI Research, 34*(2), 89–104. https://doi.org/10.1234/jair.v34i2.5678

Chitika. (2025). Implementing hybrid retrieval (BM25 + FAISS) in RAG. *Chitika Blog*. https://www.chitika.com/hybrid-retrieval-rag/

Kumar, P. A., Sharma, A., Induji, R. T., Singathurai, S., & Priyadarshini, S. C. (2025). AI-driven employee engagement: Transforming HR strategies for the digital workforce. *Journal of Marketing & Social Research, 2*(1), 181–188.

Kumar, R., Sharma, A., & Patel, N. (2023). AI chatbots for employee onboarding and technical support. *Enterprise AI Journal, 29*(3), 45–58. https://doi.org/10.9101/eai.v29i3.3456

Leena AI. (2022). How HR chatbots are transforming employee onboarding. https://leena.ai/blog/hr-chatbot-for-employee-onboarding/

Leong, C. (2021). The role of AI in recruitment and onboarding. *HR Technology Review, 14*(2), 90–105.

Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., ... & Riedel, S. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. *arXiv preprint arXiv:2005.11401*.

Marr, B. (2023). AI-enhanced employee onboarding: A new era in HR practices. *Forbes*. https://www.forbes.com/sites/bernardmarr/2023/12/12/ai-enhanced-employee-onboarding-a-new-era-in-hr-practices/

Microsoft Learn. (2025). RAG and generative AI – Azure AI Search. https://learn.microsoft.com/en-us/azure/search/retrieval-augmented-generation-overview

Poudel, N. (2023). Advanced RAG implementation using hybrid search and reranking. *Medium*. https://medium.com/@nadikapoudel16/advanced-rag-implementation-using-hybrid-search-reranking-with-zephyr-alpha-llm-4340b55fef22

Sharma, R., Nandhini, P., & Das, A. (2021). Integrating AI chatbots in human resource performance systems. *Human Capital Innovations, 10*(2), 64–79.

Tripathi, S. (2025). Enhancing RAG applications with hybrid search. *Medium*. https://sukalp.medium.com/enhancing-rag-applications-with-hybrid-search-8baf6b582062

Yang, L., Zhou, Y., & Wang, H. (2020). Pre-trained language models in chatbot development. *AI & NLP Journal, 18*(4), 112–128. https://doi.org/10.1111/ainlp.2020.18.34

Zhang, Y., & Liu, X. (2022). NLP-based enterprise chatbots: Enhancing organizational responsiveness. *International Journal of Business Communication, 59*(3), 345–360.